

Solution Explanation

In all implementation, we employ cache-friendly single dimensional arrays of size $N * N$ to represent square matrices. In implementations which uses loop tiling, the size of submatrices must divide N .

This has some advantages:

1. All elements are in one contiguous block of memory, for example $A[1 .. N * N]$
2. Only one memory access to reach the element, e.g., $A[i * N + j]$ instead of $A[i][j]$
3. Adjacent elements for both row-wise and column-wise access offer good spatial locality

Naive non-blocking implementation

The function is `void matrix_mult_naive(matrix_t *, matrix_t *, matrix_t *, double *)`.

This function uses three nested loop to multiply matrix A by matrix B. In this naive code, we don't apply optimizations, i.e., loop tiling, temporary variables, or parallel computation, etc.

```
for (int i = 0; i < N; i++)
{
    for (int j = 0; j < N; j++)
    {
        C->mem[i * N + j] = 0.0;
        for (int k = 0; k < N; k++)
        {
            C->mem[i * N + j] += A->mem[i * N + k] * B->mem[k * N + j];
        }
    }
}
```

Blocked matrix implementation

The function is `void matrix_mult_block(matrix_t *, matrix_t *, int, matrix_t *, double *)`.

Loop Tiling The matrix is divided into $N / \text{block_size}$ square sub-matrices (called blocks). The three outer loops (`bi`, `bj`, and `bk`) browse through each block.

- Loops `bi` and `bj` select the result block in matrix C
- Loop `bk` selects blocks in A and B to compute for this block in C
- The inner loops `i`, `j`, and `k` perform actual matrix multiplication.

In addition, we use a temporary variable `sum` to reduce pointer access to `C[i][j]`.

```
for (int bi = 0; bi < N; bi += block_size)
{
    for (int bj = 0; bj < N; bj += block_size)
    {
        for (int bk = 0; bk < N; bk += block_size)
        {
            int i_end = min(bi + block_size, N);
            int j_end = min(bj + block_size, N);
            int k_end = min(bk + block_size, N);

            for (int i = bi; i < i_end; i++)
            {
                for (int j = bj; j < j_end; j++)
                {
                    register double sum = 0; // register variable to store summation
```

```

        for (int k = bk; k < k_end; k++)
    {
        sum += A->mem[i * N + k] * B->mem[k * N + j]; // loop tiling
    }
    C->mem[i * N + j] = sum;
}
}
}
}
}
```

Cache Benefits

1. Spatial locality: The block is sufficient small, so access to consecutive elements in array A is predicted by the CPU for cache optimization.
 2. Temporal locality:
 - Variable `sum` is stored in CPU register for fast access.
 - Elements from $A[i * N]$ to $A[i * N + k]$ are reused $blockSize^2$ times within j loop.

Multiplication of matrix blocks using cblas_dgemm in ijk implementation

The function of this implementation is `matrix_mult_bla`.

Instead of three inner loops i , j , and k , this implementation employs Basic Linear Algebra Subprogram (BLAS) to multiply the sub-matrices.

```

for (int bi = 0; bi < N; bi += block_size)
{
    for (int bj = 0; bj < N; bj += block_size)
    {
        for (int bk = 0; bk < N; bk += block_size)
        {
            int blk_rows = min(block_size, N - bi);
            int blk_cols = min(block_size, N - bj);
            int blk_inner = min(block_size, N - bk);

            cblas_dgemm(
                CblasRowMajor,
                CblasNoTrans,
                CblasNoTrans,
                blk_rows,
                blk_cols,
                blk_inner,
                1.0,
                &A->mem[bi * N + bk],
                N,
                &B->mem[bk * N + bj],
                N,
                1.0,
                &C->mem[bi * N + bj],
                N
            );
        }
    }
}

```

Non-blocking BLAS dgemm routine

In the function `matrix_mult_cblas`, we invoke BLAS dgemm to multiply A by B.

```
cblas_dgemm(  
    CblasRowMajor,  
    CblasNoTrans,  
    CblasNoTrans,  
    N,  
    N,  
    N,  
    1.0,  
    A->mem,  
    N,  
    B->mem,  
    N,  
    0.0,  
    C->mem,  
    N);
```

Experimental Results

Environment Specification

The experiment runs on a *Ubuntu Linux* computer of which specifications are:

Component	Specification
Operating System	Ubuntu Linux
Kernel Version	6.6.87.2
Architecture	x86_64
CPU Model	AMD Processor
CPU Count	16 cores
L1 Cache	32K
L2 Cache	4 MiB
L3 Cache	32 MiB

Estimation

We first calculate the maximum block size fitting CPU cache. Let H be the maximum block size, we have:

- One block from matrix A: H^2 elements
- One block from matrix B: H^2 elements
- One block from matrix C: H^2 elements

In total, there are $3 * H^2$ elements of 64-bit floating point. Thus, CPU cache must be at least $3 * H^2 * 8$ bytes. Or, equivalently,

$$3 * H^2 * 8 \leq \text{CPU cache size}$$

$$\Leftrightarrow H \leq \sqrt{\frac{\text{CPU cache size}}{3 * 8}}$$

In our program, we only implement one layer cache (check `matrix_mult_block` and `matrix_mult_cblas_block` implementations). Hence, we only take into account the largest L3 Cache. Alternating L3 Cache size of 32 MiB = 33,554,432 bytes to the formula, we obtain $H \leq 1182.41335$ bytes. This is the block size's upper bound in the benchmark.

Secondly, for the ease of estimation, we choose matrix sizes which are multiples of 512. All matrix sizes we run our tests: 1024, 1536, 2048, 2560, 3072, 3584, and 4096. Lastly, for blocking implementations, the chosen block sizes must be common divisors of all matrix sizes.

Test Scenarios

Basing on the estimation, we provide the summary table for all test scenarios:

Implementation	Matrix Size	Block Size
Naive	1024, 1536, 2048, 2560, 3072, 3584, 4096	
Block	1024, 1536, 2048, 2560, 3072, 3584, 4096	16, 128, 256, 512, 1024
CBlas	1024, 1536, 2048, 2560, 3072, 3584, 4096	
CBlas + Block	1024, 1536, 2048, 2560, 3072, 3584, 4096	16, 128, 256, 512, 1024

Benchmark Results and Analysis

Naive and Block Implementation Comparison First we look at the runtime for two blocking implementations: `matrix_mult_block` (block) versus `matrix_mult_naive` (naive):

Matrix Size	Naive (seconds)	Best Block Block Size	Best Block Time (seconds)	Speedup (Naive/Block)
1024	9.343332	128	3.638008	2.57x
1536	27.428002	256	12.745546	2.15x
2048	69.980651	256	30.772002	2.27x
2560	140.638135	1024	60.225147	2.34x
3072	255.616116	256	107.104698	2.39x
3584	434.211685	1024	168.255812	2.58x
4096	621.830439	128	232.936458	2.67x

Key Observations:

- Loop tiling technique brings about more than two time (e.g., ~2.2 to 2.7 times) faster calculation time. The implementation deliberately omits advanced optimization such as multi-threading, kij loop ordering, SIMD, iteration unrolling, etc. This means the significant results is because of the increment in cache hit rate.

Cblas and Cblas-block Implementation Comparison Let's look at the runtime results:

Matrix Size	Cblas (seconds)	Best Cblas-block Block Size	Best Cblas-block Time (seconds)	Speedup (Cblas/Cblas-block)
1024	0.016634	1024	0.015478	1.07x
1536	0.035903	1024	0.059824	0.60x
2048	0.071432	1024	0.095189	0.75x
2560	0.140348	512	0.193757	0.72x
3072	0.218941	1024	0.295758	0.74x
3584	0.367671	1024	0.525116	0.70x
4096	0.519840	1024	0.800810	0.65x

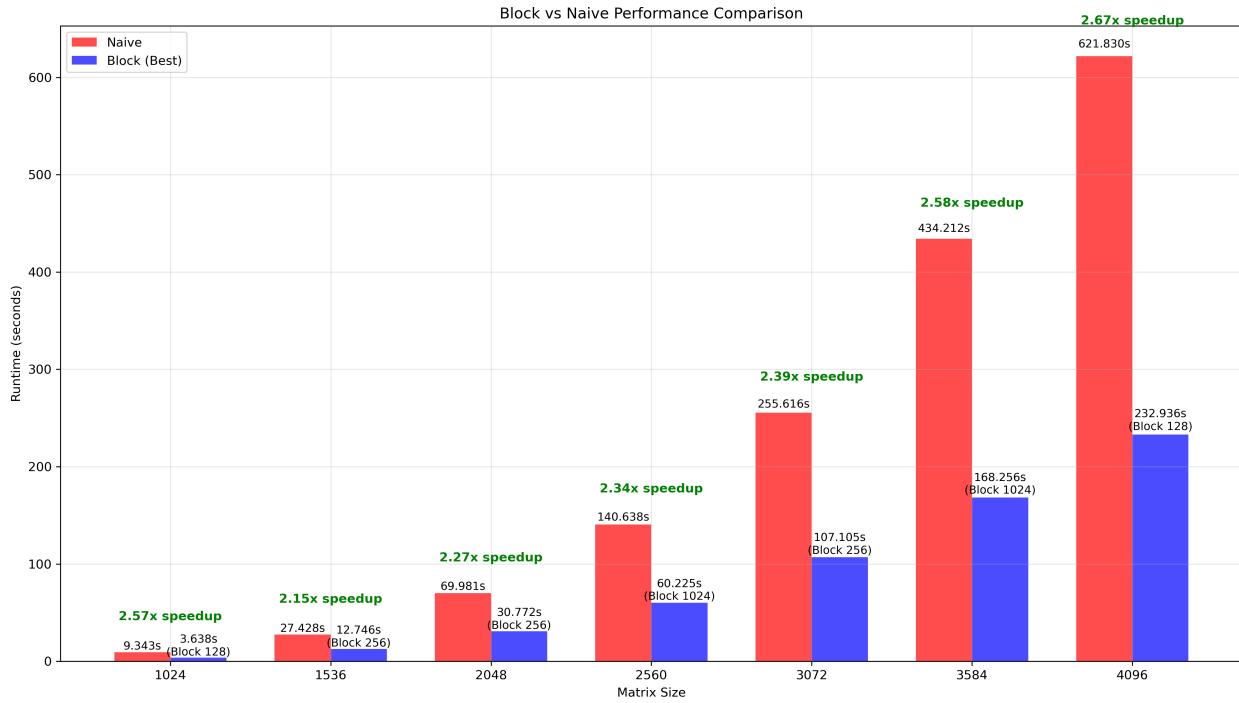


Figure 1: Naive versus Block Implementation

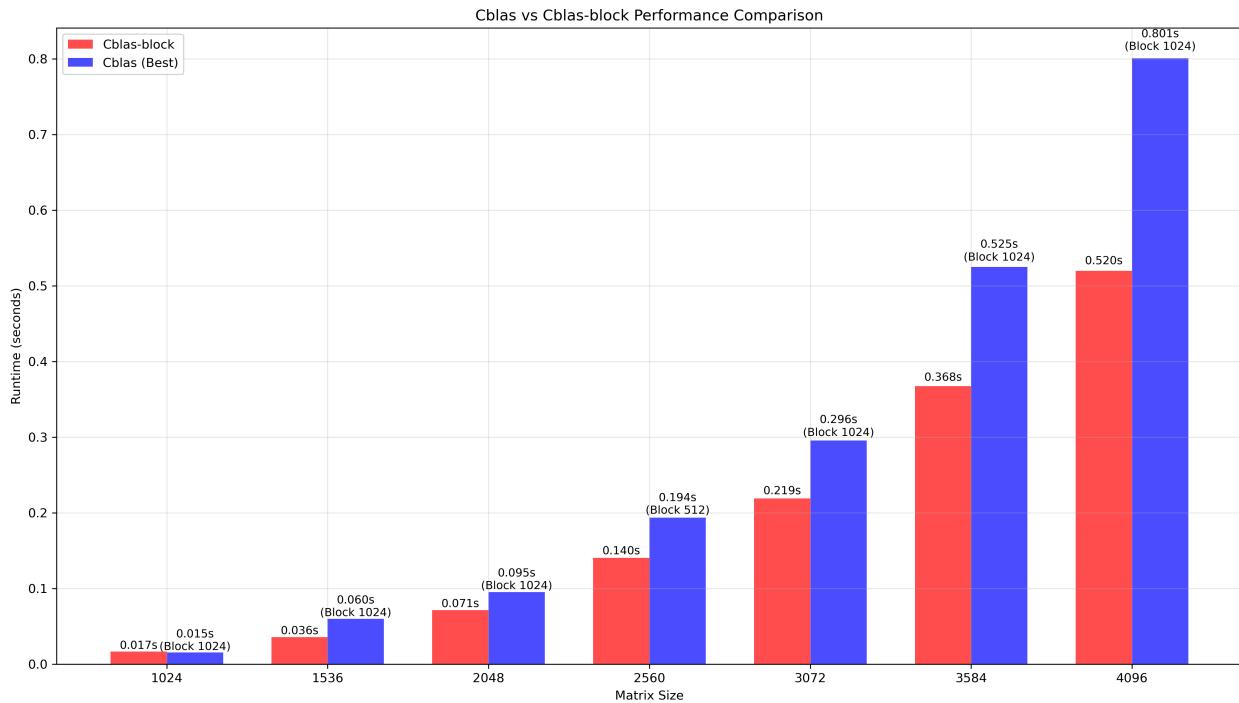


Figure 2: Cblas and Cblas-block Implementation

Key Observation:

- The loop tiling technique in `matrix_mult_blas_block` doesn't make the calculation faster than `matrix_mult_cblas`. This is because of the overheads of the three four ijk loops.