

Solution Explaination

Serial Implementation

Norm Calculation

The 1-norm of a matrix is defined as:

$$\|A\|_1 = \max_{1 \leq i \leq n} \sum_{j=1}^n |a_{ij}|$$

The function `matrix_norm_serial` process each row entirely and then update the result to the global maximum. Loop tiling is applied to optimize CPU cache.

```
for (i = 0; i < N; i++)
{
    row = &data[i * N];
    row_sum = 0.0;
    for (bj = 0; bj < N; bj += block_size)
    {
        j_end = MIN(bj + block_size, N);
        for (j = bj; j < j_end; j++)
        {
            row_sum += fabsl(row[j]);
        }
    }
    max_row_sum = MAX(max_row_sum, row_sum);
}
```

Matrix Multiplication

The function implementing the serial algorithm is `matrix_mult_serial(size_t, matrix_t *, matrix_t *, matrix_t *)`. The algorithm includes the three nested four `kij` loop. It runs in $O(n^3)$ time-complexity.

Four optimization techniques are applied to the serial implementation:

Loop tiling The three outer loops of `bi`, `bj`, and `bk` browse through $\lceil \frac{\text{Matrix Size}}{\text{Block size}} \rceil$ blocks. The block size is optimially selected to maximize cache hits.

```
for (size_t bi = 0; bi < N; bi += block_size)
{
    for (size_t bj = 0; bj < N; bj += block_size)
    {
        for (size_t bk = 0; bk < N; bk += block_size)
        {
            /*
        }
```

```

    ... Inner implementation (see kij ordering) ...
*/
```

}

}

}

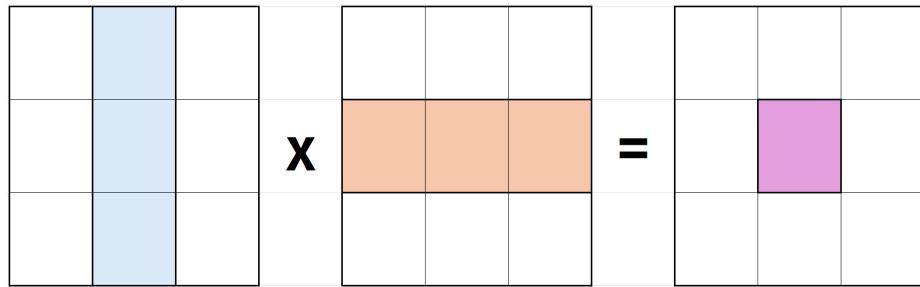


Figure 1: Block Selection

kij ordering The simplified kij implementation is as below:

```

for (k = 0; k < n; k++) {
    for (i = 0; i < n; i++) {
        int x = A[i][k];
        for (j = 0; j < n; j++) {
            C[i][j] = x * B[k][j];
        }
    }
}
```

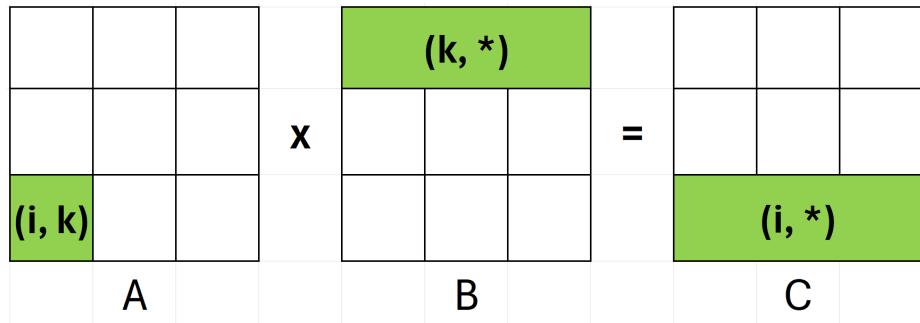


Figure 2: kij ordering

In the inner most loop j , every elements in $C[i, *]$ and $B[k, *]$ are accessed `block_size` times. The temporary locality is improved by allowing CPU to efficiently cache row $C[i]$ and row $B[k]$. To avoid multiple pointer access

to $A[i][k]$, we store $A[i][k]$ value to a temporary variable x . This technique enhances spatial locality and has cache miss rate in B and C of $\frac{1}{\text{block size}}$.

The detailed implementation is provided as below:

```

const size_t i_end = MIN(bi + block_size, N);
const size_t j_end = MIN(bj + block_size, N);
const size_t k_end = MIN(bk + block_size, N);

for (size_t k = bk; k < k_end; k++)
{
    for (size_t i = bi; i < i_end; i++)
    {
        register const double lhs_data = lhs->data[i * N + k];
        register double *rhs_row = &rhs->data[k * N];
        register double *result_row = &result->data[i * N];
        register const size_t limit = j_end - ((j_end - bj) % 4);
        register size_t j;

        for (j = bj; j < limit; j += 4)
        {
            result_row[j] += lhs_data * rhs_row[j];
            result_row[j + 1] += lhs_data * rhs_row[j + 1];
            result_row[j + 2] += lhs_data * rhs_row[j + 2];
            result_row[j + 3] += lhs_data * rhs_row[j + 3];
        }

        for (; j < j_end; j++)
        {
            result_row[j] += lhs_data * rhs_row[j];
        }
    }
}

```

Register variables

```

register const double lhs_data = lhs->data[i * N + k];
register double *rhs_row = &rhs->data[k * N];
register double *result_row = &result->data[i * N];
register const size_t limit = j_end - ((j_end - bj) % 4);
register size_t j;

```

We store critical variables for the multiplication in register for fast access.

Variable	Description
<code>lhs_data</code>	The fixed cell in matrix $A[i][j]$. <code>lhs</code> stands for left hand side operands corresponding to matrix A in the operation $A * B$
<code>rhs_row</code>	The k^{th} row of matrix B . <code>rhs</code> stands for right hand side.
<code>result_row</code>	The i^{th} row of result matrix C

Iteration unrolling We expand a loop to do multiple operations in each iteration. This improves performance by reducing the overhead of loop control. The initial loop is

```
for (j = bj; j < j_end; j++)
{
    result_row[j] += lhs_data * rhs_row[j];
}
```

By applying iteration unrolling, we process 4 elements per iteration:

```
for (j = bj; j < limit; j += 4)
{
    result_row[j] += lhs_data * rhs_row[j];
    result_row[j + 1] += lhs_data * rhs_row[j + 1];
    result_row[j + 2] += lhs_data * rhs_row[j + 2];
    result_row[j + 3] += lhs_data * rhs_row[j + 3];
}

for (; j < j_end; j++)
{
    result_row[j] += lhs_data * rhs_row[j];
}
```

Threaded Implementation

Norm Calculation

`matrix_norm_threaded` function sets up and manages the parallel norm computation:

1. Work distribution

```
const size_t N = mat->size;
const size_t PARTITION_SIZE = (size_t)ceil((double)N / (double)num_threads);
```

The matrix is horizontally divided by rows:

- Each thread handles `PARTITION_SIZE` consecutive rows

- The last thread may process fewer rows if N is not evenly divisible

2. Thread parameter initiation

```
for (size_t i = 0; i < num_threads; i++)
{
    worker_params[i].mat = mat;
    worker_params[i].shared_max_sum = &shared_result; // All threads share this
    worker_params[i].mutex = &mutex; // All threads share this mutex
    worker_params[i].start_index = i * PARTITION_SIZE;
    worker_params[i].end_index = MIN((i + 1) * PARTITION_SIZE, N);
    worker_params[i].block_size = block_size;

    pthread_create(
        &threads[i],
        NULL,
        matrix_norm_worker,
        (void *)&worker_params[i]);
}
```

3. Result aggregation

One shared variable is used for storing the computation result. A critical section is created at the end of `matrix_norm_worker` for the shared result update.

```
pthread_mutex_lock(worker_params->mutex);
*(worker_params->shared_max_sum) = MAX(*(worker_params->shared_max_sum), local_max_sum);
pthread_mutex_unlock(worker_params->mutex);
```

Matrix Multiplication

The main function that is responsible for parallel matrix multiplication is `matrix_mult_threaded`. The function has 5 main steps:

1. Setup thread and allocate memory

```
const size_t N = lhs->size;
const size_t PARTITION_SIZE = (size_t)ceil((double)N / (double)num_threads);
pthread_t *threads;
matrix_mult_worker_params_t *worker_params;

threads = (pthread_t *)calloc(num_threads, sizeof(pthread_t));
worker_params = (matrix_norm_worker_params_t *)calloc(num_threads, sizeof(matrix_norm_worker_params_t));
```

This code estimates the number of rows each thread will handle `PARTITION_SIZE` and allocates the arrays of work threads and their parameters.

2. Work distribution

```
worker_params[i].block_start_index = i * PARTITION_SIZE;
worker_params[i].block_end_index = MIN((i + 1) * PARTITION_SIZE, N);
```

The matrix is divided among threads horizontally, for example:

- Thread 0: rows 0 to PARTITION_SIZE - 1
 - Thread 1: rows PARTITION_SIZE to 2 * PARTITION_SIZE - 1
 - Thread i-th: rows $i * \text{PARTITION_SIZE}$ to $(i + 1) * \text{PARTITION_SIZE} - 1$
3. Separate result buffers

```
worker_params[i].result = (double *)calloc(N * N, sizeof(double));
```

Each thread has a complete individual NxN result matrix. This avoid race-condition without using locking mechanism, e.g., mutexes, which serverely drags the performance down due to overheads.

4. Thread creation and execution

```
pthread_create(&threads[i], NULL, matrix_mult_worker, (void *)&worker_params[i]);
```

5. Result aggregation

```
for (size_t i = 0; i < num_threads; i++)  
{  
    pthread_join(threads[i], NULL);  
    for (size_t j = 0; j < N * N; j++)  
    {  
        result->data[j] += worker_params[i].result[j];  
    }  
    free(worker_params[i].result);  
}
```

Experiment Results

The system configurations in which the tests are performed are:

Component	Specification
Operating System	Ubuntu Linux
Kernel Version	6.6.87.2
Architecture	x86_64
CPU Model	AMD Processor
CPU Count	16 cores
L1d Cache	32K
L2 Cache	4 MiB
L3 Cache	32 MiB

Based on the system information, we conducted the experiments with the matrix size of 512-multiples. The block size is 512 to fit the CPU caches and divides

all matrix sizes. Second, the number of threads in the parallel test matches the CPU core number, that is, 16.

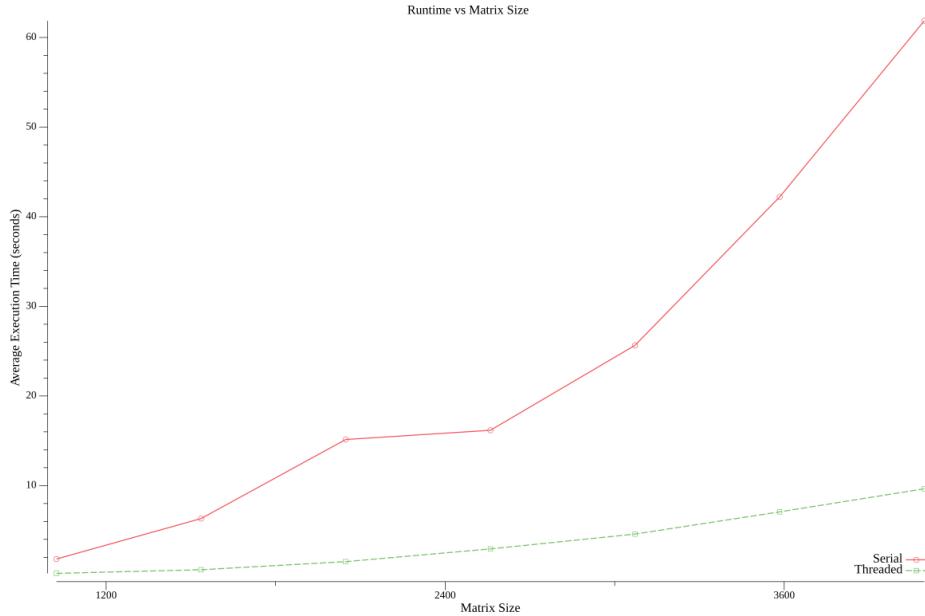
Each test scenario is run fifty times and the runtime of each run is recorded for average estimation.

Test Type	Matrix Size	Block Size	Number of Threads	Number of Iterations	Implementation
Serial	1024	512	1	50	serial
Serial	1536	512	1	50	serial
Serial	2048	512	1	50	serial
Serial	2560	512	1	50	serial
Serial	3072	512	1	50	serial
Serial	3584	512	1	50	serial
Serial	4096	512	1	50	serial
Threaded	1024	512	16	50	threaded
Threaded	1536	512	16	50	threaded
Threaded	2048	512	16	50	threaded
Threaded	2560	512	16	50	threaded
Threaded	3072	512	16	50	threaded
Threaded	3584	512	16	50	threaded
Threaded	4096	512	16	50	threaded

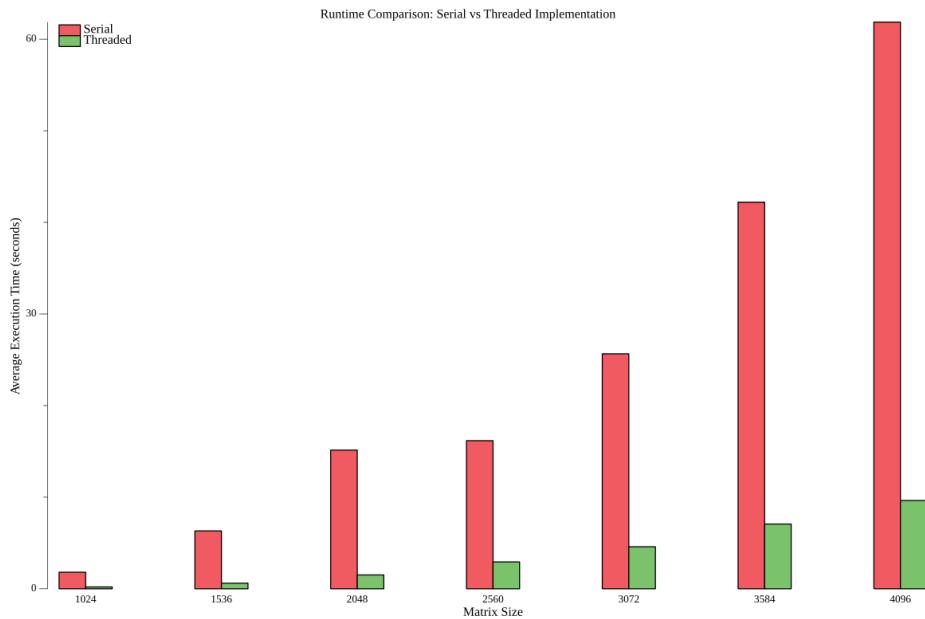
The experimental results table is:

Matrix Size	Serial Time (s)	Threaded Time (s)	<u>Threaded Time / Serial Time</u>
1024	1.809551	0.206279	8.77
1536	6.319773	0.611323	10.34
2048	15.141643	1.514556	10.00
2560	16.170770	2.935385	5.51
3072	25.652109	4.590539	5.59
3584	42.207631	7.070302	5.97
4096	61.862645	9.640536	6.42

The Dependency of Matrix Size and Runtime



The Speedup



Key Observations:

1. Decreasing speedup with matrix size
 - With small matrices (1024 - 2048), from 8.77x to 10.34x speedup
 - With large matrices (2560 - 4096), from 5.51x to 6.42x speedup
 - Peak performance is at 1536x1536 matrices with 10.34x speedup

This is caused by memory pressure and overheads which come along with larger matrices. Performance reaches its peak at 1536x1536 and 2048x2048 matrices, when there is good balance between parallelization and cache fitness.

2. Inspite of 16 CPU cores, the speedup reaches only ~10x (not 16x). As the matrix size grows larger, the speedup decreases significantly.

All 16 cores compete for the same memory bus, causing memory bandwidth bottlenecks.

References

- Cache Effect in Matrix Multiplication. Tutorial for Assignment 1
- Wikipedia - Loop unrolling